

Names: Yen-Wen Lu, Ming-Cheng Liao, Yu-Tang Su

Netid: ywlu2, mcliao2, ytsu2

rai ids: 5d97b1e588a5ec28f9cb9493

5d97b1df88a5ec28f9cb9488

5d97b21188a5ec28f9cb94e5

Team name: dpp

School affiliation: on-campus

1. Optimization Approach and Results

- **Optimization 4 (in Line 559 in new-forward.cuh): Unroll + shared-memory Matrix multiply + Exploiting parallelism in input images, input channels, and output channels.**

- **how you identified the optimization opportunity**

In milestone2, Matrix multiplication was implemented with three four loops, which indicated that lots of multiplication was not paralleled. Therefore, we think that Unroll X is able to reduce for loops and then accelerate the computation. Moreover, we tried to launch a kernel with two dimensional grid (X dimension is the size of unroll x and Y dimension is the size of batch) to calculate unroll X with all batch simultaneously, which is also the exploiting parallelism in input images optimization. We think it would faster than using for loop to compute each images in each batch. In addition, in milestone2, we accessed data from global memory which was slow. It would be faster if we accessed data from shared memory when doing matrix multiplication.

- **Why do you thought the approach would be fruitful**

Unrolling X is able to effectively reduce lots of iteration in kernel function. Therefore, we believe unrolling X is a fruitful method.

Weight Matrices and unrolling X were stored in shared memories when doing matrix multiplication, which can be accessed faster by threads than those stored in global.

Exploiting parallelism in input images, input channels, and output channels is faster since threads implemented each elements wise multiplication simultaneously than using for loop of each input images, input channels, and output channels to implemented matrix multiplication.

- **Result**

Fig1. Indicates that this method is slower compared to milestone 2.1. The possible reason is that global memory is not enough to store unroll X. We had to use three kernels to load part of batch images separately. We separated 10000 batch size into 4000, 4000, 2000 to do the Unroll + shared-memory Matrix multiplication. Since we called kernel function for three times, op time will be slower than the op time in milestone 2.1. From Fig2, though kernel function run quite fast, three cudamalloc and

cudaFree operation spent about 0.4s, which is the bottle neck in all operation time and it is the main reason why optimization 4 is slower than milestone2.

```
* Running python m4.1.py 10000
Loading fashion-mnist data... done
Loading model... done
New Inference
B = 10000, C = 1, M = 12, K = 5, H_out = 66, W_out = 66
Op Time: 0.082661
B = 10000, C = 12, M = 24, K = 5, H_out = 29, W_out = 29
Op Time: 0.505352
Correctness: 0.7653 Model: ece408
```

Fig 1. Op time of Optimization 4

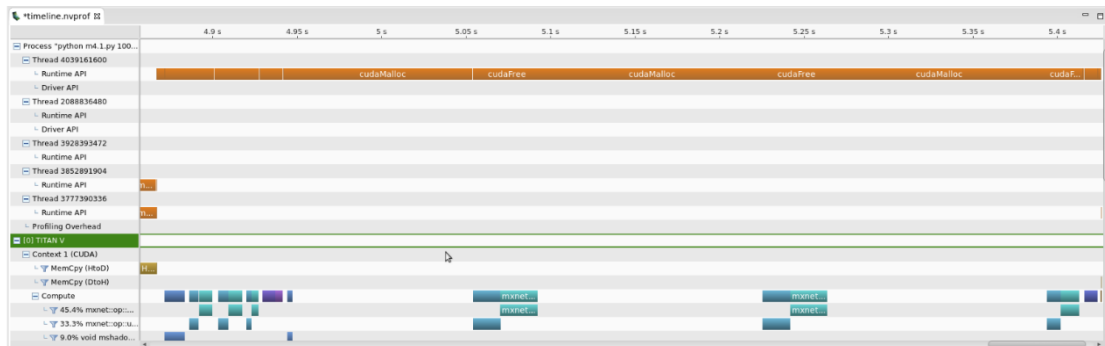


Fig 2. NVPP Profile of Optimization 4

- **Optimization 5 (in Line 627 in new-forward.cuh): Unroll + Register-Tile Matrix multiply + Exploiting parallelism in input images, input channels, and output channels.**

■ how you identified the optimization opportunity

From lecture slides of “Lecture 3: Thread Coarsening and Register Tiling” in ECE508[1], Registers do not require memory access instructions, which indicates high throughput. Moreover, this optimization is implemented with thread coarsening, so the computation from merged threads can now share registers.

■ **Why do you think the approach would be fruitful**

Since registers is not capable of all data from global memory. What we did is loading weight matrix in register and loading x into shared memory and then doing the matrix multiplication. Since registers' throughput is extremely high, this approach would be fruitful.

■ Result

In theory, register should be much faster. However, the result in Fig 3. Showed that optimization 5 is slower than Optimization 4. From Fig 4. , bottle neck is executing

kernel function, which spent longer time than Optimization 4. The possible reason for longer execution time is that the kernel spent too much time loading unroll X into shared memory, although the kernel accessed weight matrix fast form registers.

```
* Running python m4.1.py 10000
Loading fashion-mnist data... done
Loading model... done
New Inference
B = 10000, C = 1, M = 12, K = 5, H_out = 66, W_out = 66
Op Time: 0.060375
B = 10000, C = 12, M = 24, K = 5, H_out = 29, W_out = 29
Op Time: 0.437847
Correctness: 0.7653 Model: ece408
```

Fig 3. Op time of Optimization 5

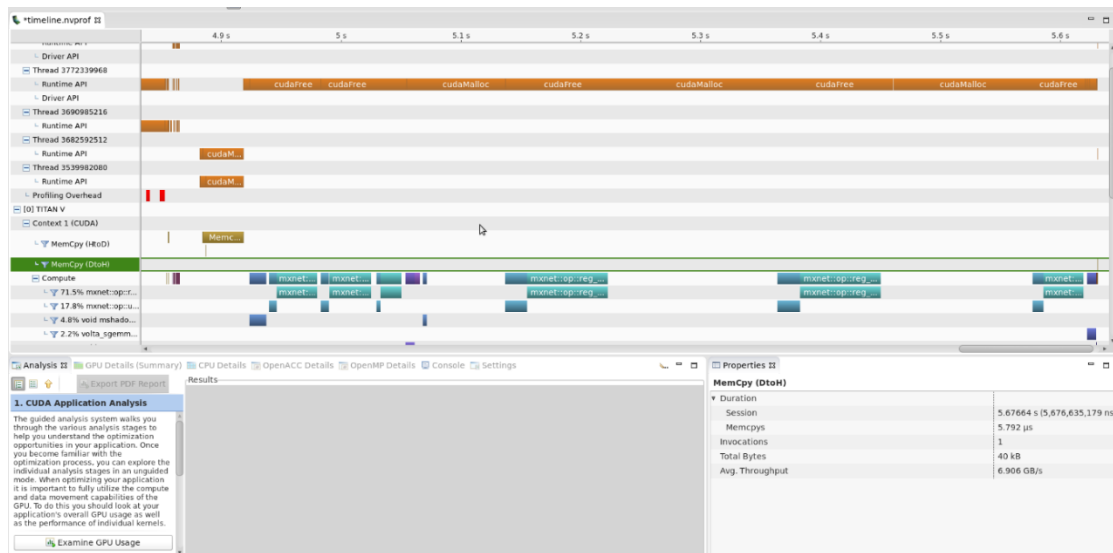


Fig 4. NVPP Profile of Optimization 5

- **Optimization 6 (in Line 689 in new-forward.cuh): Kernel fusion for unrolling and matrix-multiplication (unrolling X into shared memory directly)**

■ How you identified the optimize opportunity

We thought that combining unrolling X kernel and matrix multiplication into one kernel would be faster since this only did memcopy from host and device and memcopy from device to host one time. Moreover, we were able to do unrolling X, store part of X into shared memory and matrix multiplication in a loop.

■ Why do you think the approach would be fruitful?

Kernel fusion combines two kernels (unrolling and matrix-multiplication) into one kernel. This method only takes one time to allocate memory on device and announce

Fig5. indicates that this kernel fusion method is faster than milestone 2. There are two reasons why this optimization method is much faster. First, we loaded part of data into data for one time and compute part of elements, so there will not be insufficient memory problem. Second, our kernel fusion implemented unrolling and matrix multiplication in the same loop, which is faster than separating unrolling X and matrix multiplication.

```
.nvprof python m4.1.py 10000
Loading fashion-mnist data... done
==452== NVPROF is profiling process 452, command: python m4.1.py 10000
Loading model... done
New Inference
Op Time: 0.019478
Op Time: 0.063723
Correctness: 0.7653 Model: ece408
```

The screenshot displays the NVIDIA Nsight Systems performance analysis tool. The top timeline shows the execution of a CUDA application, with various kernels and streams color-coded. The bottom panel displays the 'Properties' tab for the 'MemCpy (HtoD)' kernel, showing its duration, session, and memory usage.

Properties (HtoD)

- Duration: 4.99147 s (4,991,469,788 ns)
- Session: 35,834,781 ms
- MemCpy: 20
- Invocations: 206,996 MB
- Total Bytes: 5,832 GB/s
- Avg. Throughput: 5,832 GB/s

Fig 6. NVPP Profile of Optimization

- **Optimization 7 (in Line 507 in new-forward.cuh):** Sweeping various parameters to find best values (Optimization 7 is the best method)

■ How you identified the optimization opportunity

Certain block sizes or other parameters may avoid control divergence problem or memory not coalescing. Moreover, if setting the block size too large, it may cause illegal memory access problem. Therefore, sweeping various parameters to find best values is able to prevent above problems to occur and improve computation speed.

We found that Op1 and Op2 s' best block size for Op time are different. Our method was to test each Ops' best block size. Since Op1 is for input channel equaling 1 and Op2 is for input channel equaling 12. We used if statement to control different block size for different Op. Furthermore, we found that using __restrict__ to one parameter for the whole kernel function is better for optimization than using __restrict__ to all parameters. That is why we think this approach is fruitful.

■ Result

Table 2. shows that our all testing with different block size and optimization method. From the "Total Op time" we can see that kernel fusion method with Block size equaling 24 and kernel fusion with Block equaling 32 had the shortest Op time. However, from Fig 9. And Fig 10. , which is divergent branches in nvpp profile, it shows that our code exists control divergence. Therefore, Op time may be shorter if we could solve this control divergence problem.

	Optimization Method(Op1)	Optimization Method(Op2)	Block Size(Op1)	Block Size(Op2)	Total Op Time(sec)
Test 1	Baseline method	Baseline method	16	16	0.115309
Test 2	Baseline method	Baseline method	16	32	0.085582
Test 3	Fusion method	Fusion Method	16	32	0.087712
Test 4	Fusion method	Fusion method	6	6	0.223641
Test 5	Fusion method	Fusion method	12	12	0.091421
Test 6	Fusion method	Fusion method	18	18	0.104445
Test 7	Fusion method	Fusion method	24	24	0.089376
Test 8	Fusion method	Fusion method	32	32	0.105804
Test 9	Fusion method	Fusion method	16	32	0.078843
Test 10	Fusion method	Fusion method	16	24	0.065626

Table 2. Different parameters test

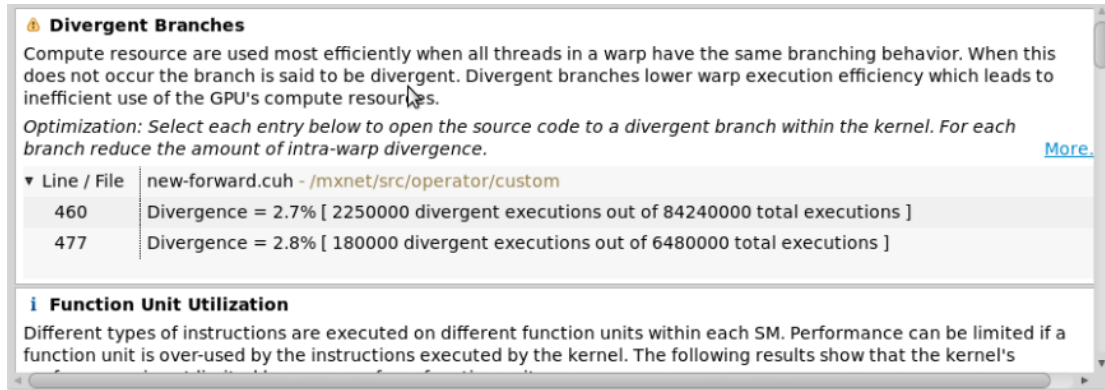


Fig 10. Divergent Branches for Op2

1.1 implemented Optimization

Table 2. shows which optimization methods listed on github were implemented in our optimization kernels.

Table 3. marks each optimization's position in new-forward.cuh.

Optimization	Optimization 4	Optimization 5	Optimization 6	Optimization 7
Optimization method in ECE 408 project github				
Unroll + shared-memory Matrix multiply	V			
Kernel fusion for unrolling and matrix-multiplication			V	V
register-tiled		V		
Sweeping various parameters to find best values				V
Exploiting parallelism in input images, input channels, and output channels	V	V	V	V

Table 2. Optimization List

Optimization	Optimization 4	Optimization 5	Optimization 6	Optimization 7
Line number in new-forward.cuh	559	627	689	507

Table 3. Line number of each Optimization in new-forward.cuh

2. References (as needed)

[1] W. Hwu, "Lecture 4: Joint Register and Shared-Memory Tiling" ECE 508 Lecture Slides, 2019.

[2] W. Hwu, "Application Case Study –Machine Learning" ECE 408 Lecture Slides, 2019.

3. Contribution: How your team organized and divided up this work

Team member	Contribution
Ming-Cheng Liao	Optimization 4, Optimization 5, Optimization 6, Optimization 7
Yen-Wen Lu	Optimization 4, Optimization 6, Optimization 7, Final Report
Yu-Tang Su	Optimization 4, Optimization 7

4. Previous Report

1. Include a list of all kernels that collectively consume more than 90% of the program time

CUDA memcpy HtoD

volta_scudnn_128x64_relu_interior_nn_v1

volta_gcgemm_64x32_nt

fft2d_c2r_32x32

volta_sgemm_128x128_tn

op_generic_tensor_kernel

fft2d_r2c_32x32

2. Include a list of all CUDA API calls that collectively consume more than 90% of the program time.

cudaMemGetInfo

cudaStreamCreateWithFlags

cudaFree

3. Explanation of the difference between kernels and API calls

Kernels are the function that executed on GPU. API calls are the functions used to interact between host and device.

4. Show output of rai running MXNet on the CPU

Loading fashion-mnist data... done

Loading model... done

New Inference

EvalMetric: {'accuracy': 0.8154}

5. List program run time

4.97s

6. Show output of rai running MXNet on the GPU

Loading fashion-mnist data... done

Loading model... done

New Inference

EvalMetric: {'accuracy': 0.8154}

7. List program run time

4.66s

8. Create a CPU implementation

```
void forward(mshadow::Tensor<cpu, 4, DType> &y, const
mshadow::Tensor<cpu, 4, DType> &x, const
mshadow::Tensor<cpu, 4, DType> &k)
{
    /*
        Modify this function to implement the forward pass
        described in Chapter 16.
        The code in 16 is for a single image.
        We have added an additional dimension to the tensors
        to support an entire mini-batch
        The goal here is to be correct, not fast (this is the
        CPU implementation.)
    */

    const int B = x.shape_[0];
    const int M = y.shape_[1];
    const int C = x.shape_[1];
    const int H = x.shape_[2];
    const int W = x.shape_[3];
```

```

const int K = k.shape_[3];

for (int b = 0; b < B; ++b) {
    for(int m = 0; m < M; ++m){
        for(int h = 0; h < H; ++h){
            for(int w = 0; w < W; ++w){
                y[b][m][h][w] = 0;
                for(int c = 0; c < C; ++c){
                    for(int p = 0; p < K; ++p){
                        for(int q = 0; q < K; ++q){
                            y[b][m][h][w] +=
x[b][c][h + p][w + q] * k[m][c][p][q];
                        }
                    }
                }
            }
        }
    }
}

```

9. List whole program execution time

m1.1.py:

user time = 24.66

system time = 5.48

elapsed time = 0:11.35

m1.2.py:

user time = 4.99

system time = 3.34

elapsed time = 0:04.73

m2.1.py:

user time = 105.00

system time = 10.65

elapsed time = 1:33.88

10. List Op Times

optime1: 11s

optime2: 78s

Implement a GPU Convolution

```

#ifndef MXNET_OPERATOR_NEW_FORWARD_CUH_
#define MXNET_OPERATOR_NEW_FORWARD_CUH_

#include <mxnet/base.h>

# define TILE_WIDTH 16
namespace mxnet
{
namespace op
{

__global__ void forward_kernel(float *y, const float *x, const float *k, const int B,
const int M, const int C, const int H, const int W, const int K)
{

    /*
    Modify this function to implement the forward pass described in Chapter 16.
    We have added an additional dimension to the tensors to support an entire mini-
    batch
    The goal here is to be correct AND fast.
    We have some nice #defs for you below to simplify indexing. Feel free to use
    them, or create your own.
    */

    const int H_out = H - K + 1;
    const int W_out = W - K + 1;
    (void)H_out; // silence declared but never referenced warning. remove this line
    when you start working
    (void)W_out; // silence declared but never referenced warning. remove this line
    when you start working

    // An example use of these macros:
    // float a = y4d(0,0,0,0)
    // y4d(0,0,0,0) = a
    #define y4d(i3, i2, i1, i0) y[(i3) * (M * H_out * W_out) + (i2) * (H_out * W_out) +
    (i1) * (W_out) + i0]
    #define x4d(i3, i2, i1, i0) x[(i3) * (C * H * W) + (i2) * (H * W) + (i1) * (W) + i0]

```

```

#define k4d(i3, i2, i1, i0) k[(i3) * (C * K * K) + (i2) * (K * K) + (i1) * (K) + i0]
    int W_grid = ceil((float)W_out / TILE_WIDTH);
    const int b = blockIdx.x;
    const int m = blockIdx.y;
    const int h = blockIdx.z / W_grid * TILE_WIDTH + threadIdx.y;
    const int w = blockIdx.z % W_grid * TILE_WIDTH + threadIdx.x;
    if(h < H_out && w < W_out){
        float val = 0;
        for(int c = 0; c < C; c++){
            for(int p = 0; p < K; p++){
                for(int q = 0; q < K; q++){
                    val += x4d(b, c, h+p, w+q) * k4d(m, c, p, q);
                }
            }
        }
        y4d(b, m, h, w) = val;
    }

```

```

#undef y4d
#undef x4d
#undef k4d
}

```

```

/*

```

This function is called by new-inl.h

Any code you write should be executed by this function.

For ECE408, we only expect the float version of the operator to be called, so here we specialize with only floats.

```

*/

```

```

template <>

```

```

void forward<gpu, float>(mshadow::Tensor<gpu, 4, float> &y, const
mshadow::Tensor<gpu, 4, float> &x, const mshadow::Tensor<gpu, 4, float> &w)
{

```

```

    // Use mxnet's CHECK_EQ to do assertions.

```

```

    // Remove this assertion when you do your implementation!

```

```

//CHECK_EQ(0, 1) << "Remove this line and replace with your
implementation";

// Extract the tensor dimensions into B,M,C,H,W,K
// ...

const int B = x.shape_[0];
const int C = x.shape_[1];
const int H = x.shape_[2];
const int W = x.shape_[3];

const int M = y.shape_[1];
const int K = w.shape_[3];

int H_out = H - K + 1;
int W_out = W - K + 1;

int W_grid = ceil(1.0*W_out / TILE_WIDTH);
int H_grid = ceil(1.0*H_out / TILE_WIDTH);

int Z = H_grid * W_grid;

// Set the kernel dimensions
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1);
dim3 gridDim(B, M, Z);

// Call the kernel
//forward_kernel<<<gridDim,
blockDim,0,s>>>(y.dptr_,x.dptr_,w.dptr_,B,M,C,H,W,K);
    forward_kernel<<<gridDim,
blockDim,0,0>>>(y.dptr_,x.dptr_,w.dptr_,B,M,C,H,W,K);
    // Use MSHADOW_CUDA_CALL to check for CUDA runtime errors.
    MSHADOW_CUDA_CALL(cudaDeviceSynchronize());

}

```

```

/*
    This tells mxnet how to do an op when it's not a float.
    This is not used in the ECE408 project
*/
template <typename gpu, typename DType>
void forward(mshadow::Tensor<gpu, 4, DType> &y, const mshadow::Tensor<gpu, 4,
DType> &x, const mshadow::Tensor<gpu, 4, DType> &w)
{
    //CHECK_EQ(0,1) << "Remove this line and replace it with your
implementation.";
}
}
}

#endif

```

Correctness and timing with 3 different dataset sizes

*** Running python m3.1.py 100**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.000270

Op Time: 0.000931

Correctness: 0.76 Model: ece408

*** Running python m3.1.py 1000**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.002707

Op Time: 0.009456

Correctness: 0.767 Model: ece408b

*** Running python m3.1.py 10000**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.028302

Op Time: 0.087000

Correctness: 0.7653 Model: ece408

***** Running nvprof python m3.1.py

Loading fashion-mnist data... done

==265== NVPROF is profiling process 265, command: python m3.1.py

Loading model... done

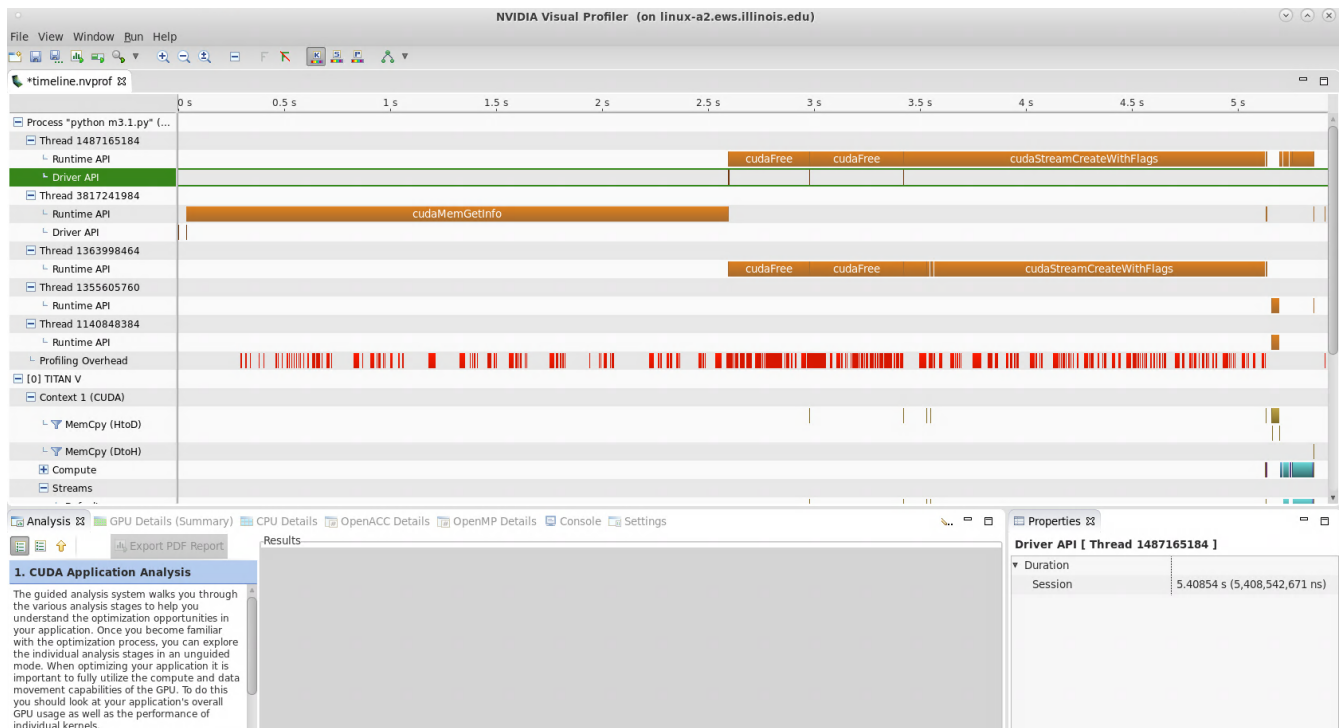
New Inference

Op Time: 0.032408

Op Time: 0.097270

Correctness: 0.7653 Model: ece408

Demonstrate nvprof profiling the execution



From this picture, we can directly see the timeline of each api like cudaFree and cudaMemGetInfo. Therefore, we can easily understand which part is the bottleneck of the process. In this case, it is cudaMemGetInfo.

Milestone4:

Implement final GPU optimizations

Optimization our team used :

- Shared Memory convolution
- Weight matrix (kernel values) in constant memory
- Tuning with restrict and loop unrolling

Implementation (we combined these 3 method into one kernel function and used new-forward.cuh):

```
ifndef MXNET_OPERATOR_NEW_FORWARD_CUH_
#define MXNET_OPERATOR_NEW_FORWARD_CUH_
```



```
#include <mxnet/base.h>
```

```
namespace mxnet
```

```
{
```

```
namespace op
```

```
{
```

```
#define y4d(i3, i2, i1, i0) y[(i3) * (M * H_out * W_out) + (i2) * (H_out * W_out) +  
(i1) * (W_out) + i0]
```

```
#define x4d(i3, i2, i1, i0) x[(i3) * (C * H * W) + (i2) * (H * W) + (i1) * (W) + i0]
```

```
#define k4d(i3, i2, i1, i0) kernel[(i3) * (C * K * K) + (i2) * (K * K) + (i1) * (K) + i0]
```

```
//#define MIN( a, b ) ( (a < b) ? a : b )
```

```
constexpr int TILE_WIDTH = 8;
```

```
__constant__ float kernel[14112];
```

```
__global__ void forward_kernel(float * __restrict__ y, const float * __restrict__ x,  
const int B, const int M, const int C, const int H, const int W, const int K, int W_grid)  
{
```

```
// use stratege 3
```

```
__shared__ float x_tile[TILE_WIDTH][TILE_WIDTH];
```

```
int H_out = H - K + 1;
```

```
int W_out = W - K + 1;
```

```
int tx = threadIdx.x; int ty = threadIdx.y;
```

```
int bx = blockIdx.x; int by = blockIdx.y;
```

```
int w = blockIdx.z % W_grid * TILE_WIDTH + tx;
```

```
int h = blockIdx.z / W_grid * TILE_WIDTH + ty;
```

```
int tile_begin_x = blockIdx.z % W_grid * TILE_WIDTH;
```

```
int tile_end_x = tile_begin_x + TILE_WIDTH;
```

```
int tile_begin_y = blockIdx.z / W_grid * TILE_WIDTH;
```

```

int tile_end_y = tile_begin_y + TILE_WIDTH;

int shift = K/2;

float val = 0;
#pragma unroll 2
for (int c = 0; c < C; c++) {

    //Loading core part of input into shared matrix
    if (w < W_out && h < H_out) {
        x_tile[ty][tx] = x4d(bx, c, h + shift, w + shift);
    } else {
        x_tile[ty][tx] = 0;
    }
    __syncthreads();

    //Set Boundary Condition
    int input_begin_x = tile_begin_x + shift;
    int input_end_x = tile_end_x + shift;
    int input_begin_y = tile_begin_y + shift;
    int input_end_y = tile_end_y + shift;

    for (int i = 0; i < K; i++) {
        for (int j = 0; j < K; j++) {
            int input_glob_x = ((w-shift)+shift) + i; //first shifting to input
            coord., and then staring convolution from the left top corner
            int input_glob_y = ((h-shift)+shift) + j; //first shifting to input
            coord., and then staring convolution from the left top corner
            int input_share_x = tx+i-shift;
            int input_share_y = ty+j-shift;
            if (input_glob_x < W && input_glob_y < H) {
                int x_min = ( (input_end_x < W_out + shift) ? input_end_x :
W_out + shift );
                int y_min = ( (input_end_y < W_out + shift) ? input_end_y :
W_out + shift );

                //using shared memory data

```

```

        if (input_glob_x >= input_begin_x && input_glob_x <
x_min && input_glob_y >= input_begin_y && input_glob_y < y_min) {
            val += x_tile[input_share_y][input_share_x] * k4d(by,
c, j, i);
        }
        else {
            //using global memory data
            val += x4d(bx, c, input_glob_y, input_glob_x) *
k4d(by, c, j, i);
        }
    }
}
__syncthreads(); // make sure all threads finish its computation before
loading new data into shared memory
}

if (w < W_out && h < H_out) {
    y4d(bx, by, h, w) = val;
}
}

```

/*

This function is called by new-inl.h

Any code you write should be executed by this function.

For ECE408, we only expect the float version of the operator to be called, so here we specialize with only floats.

*/

template <>

```

void forward<gpu, float>(mshadow::Tensor<gpu, 4, float> &y, const
mshadow::Tensor<gpu, 4, float> &x, const mshadow::Tensor<gpu, 4, float> &w)
{

```

// Extract the tensor dimensions into B,M,C,H,W,K

const int B = x.shape_[0]; // batch size

const int C = x.shape_[1]; // input channel

const int H = x.shape_[2]; // height of the input feature map

```

const int W = x.shape_[3];    // width of the input feature map
const int M = y.shape_[1];    // output channel
const int K = w.shape_[3];    // the size of the filter (assuming squared filter)
//printf("K[0]=%d\nK[1]=%d\nK[2]=%d\nK[3]=%d\n", w.shape_[0],
w.shape_[1], w.shape_[2], w.shape_[3]);

    cudaMemcpyToSymbol(kernel, w.dptr_, sizeof(float) * w.shape_[0] *
w.shape_[1] * w.shape_[2] * w.shape_[3], 0, cudaMemcpyHostToDevice);
    // Set the kernel dimensions
    int H_out = H - K + 1;
    int W_out = W - K + 1;

    int H_grid = ceil(H_out / (1.0 * TILE_WIDTH));
    int W_grid = ceil(W_out / (1.0 * TILE_WIDTH));
    int Z = H_grid * W_grid;

    dim3 gridDim(B, M, Z);
    dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1);

    forward_kernel<<<gridDim, blockDim>>>(y.dptr_, x.dptr_, B, M, C, H, W, K,
W_grid);
    // Use MSHADOW_CUDA_CALL to check for CUDA runtime errors.
    MSHADOW_CUDA_CALL(cudaDeviceSynchronize());

}

/*
    This tells mxnet how to do an op when it's not a float.
    This is not used in the ECE408 project
*/
template <typename gpu, typename DType>
void forward(mshadow::Tensor<gpu, 4, DType> &y, const mshadow::Tensor<gpu, 4,
DType> &x, const mshadow::Tensor<gpu, 4, DType> &w)
{
    CHECK_EQ(0,1) << "Remove this line and replace it with your
implementation.";
}

```

```
}  
}  
  
#endif
```

Describe and analyze the optimizations:

We integrated three optimization into on kernel. First, we loaded kernel matrix into constant memory, and then we implemented strategy 3 for convolution. For pointer y and pointer x, we added `__restricted__` to do tuning with restrict. For for loop in the kernel function, we implemented `#pragma unroll 2` to do loop unrolling.

Correctness and timing with 3 different dataset sizes

*** Running python m4.1.py 100**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.000459

Op Time: 0.001862

Correctness: 0.76 Model: ece408

*** Running python m4.1.py 1000**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.004259

Op Time: 0.018774

Correctness: 0.767 Model: ece408

*** Running python m4.1.py 10000**

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 0.057039

Op Time: 0.190717
Correctness: 0.7653 Model: ece408

use NVVP to analyze your optimization



Results	mxnet::op::forward_kernel(float*, float const *, int, int, int, int, int, int, int)
Low Memcpy/Kernel Overlap [0 ns / 5.95696 ms = 0%] The percentage of time when memcpy is being performed in parallel with kernel is low. More	Queued n/a
Low Kernel Concurrency [0 ns / 26.60715 ms = 0%] The percentage of time when two kernels are being executed in parallel is low. More	Submitted n/a
Low Memcpy Throughput [524.518 MB/s avg, for memcpys accounting for 0.6% of all memcpy time] The memory copies are not fully using the available host to device bandwidth. More	Start 4.96202 s (4,962,020,286 ns)
Low Memcpy Overlap [0 ns / 3.552 μs = 0%] The percentage of time when two memory copies are being performed in parallel is low. More	End 4.9829 s (4,982,901,961 ns)
Low Compute Utilization [29.24165 ms / 5.0302 s = 0.6%] The multiprocessors of one or more GPUs are mostly idle. More	Duration 20.88168 ms (20,881,675 ns)
	Stream Default
	Grid Size [1000,24,16]
	Block Size [8,8,1]
	Registers/Thread 40
	Shared Memory/Block 256 B
	Launch Type Normal
	▼ Occupancy
	Theoretical 75%
	▼ Shared Memory Configuration
	Shared Memory Executed 8 KiB

From “Examine Individual Kernels” in NVPP, There is no control divergence in my code. However, memcpy spent most of time and the parallelism of memcpy is not good. This may be the reason why our team’s optimization were not up to our expectation.